

Behavioral Cloning for Autonomous Steering: A PilotNet Case Study

Munjalsinh Zala — SoSe24-ML, Exam 3, Own ML Project Report — August 2026

Abstract. This report develops, evaluates, and deploys a behavioral-cloning steering-angle predictor (PilotNet; Bojarski et al., 2016) for the Udacity Self-Driving Car Simulator. A from-scratch PilotNet trained on ~7,300 combined-track recordings is compared against a simpler baseline recipe using the identical architecture and data split; the optimized configuration — primarily a data-rebalancing strategy that treats real and synthetic training examples asymmetrically rather than just equalizing class counts — improves R^2 from 0.437 to 0.489 and direction accuracy from 51.2% to 54.6%, and drives the simulator successfully in closed loop. A side comparison against a pretrained ResNet18 backbone finds the compact, 287K-parameter, from-scratch model performs at least as well as an 11.2M-parameter pretrained alternative, evidence that model capacity is not the binding constraint on this dataset. The report concludes that the amount of genuinely diverse real training data, not the model or training recipe, is the ceiling on further improvement.

1. Introduction and Problem Relevance

Behavioral cloning is a form of imitation learning in which a control policy is learned directly from demonstrations, without an explicit model of the environment's dynamics. In the context of autonomous driving, this means learning a mapping from raw camera pixels straight to a driving command — here, a steering angle — by observing a human operator's recorded driving behavior. This approach was popularized by Bojarski et al.'s 2016 NVIDIA paper, *End to End Learning for Self-Driving Cars*, which introduced the compact convolutional architecture ("PilotNet") used as the basis for this project.

The problem is relevant on several levels. Technically, end-to-end behavioral cloning replaces a hand-engineered perception→planning→control pipeline with a single learned function, data-efficient to prototype but raising hard questions about robustness and interpretability that this report investigates directly. Commercially, vision-based lane-keeping and steering-assist features are now standard in consumer ADAS, and behavioral cloning is an accessible way to prototype such policies without hand-labeled lane or object annotations. Socially, understanding *where* and *why* such models fail matters for public trust and regulatory scrutiny of vision-based driving systems — a theme this report returns to in Sections 5 and 6.

Concretely, this project trains a model to predict a continuous steering angle from a single front-facing camera frame in a simulated driving environment (the Udacity Self-Driving Car Simulator). Using a simulator rather than a real vehicle is itself a relevant methodological choice: it isolates the perception-to-control learning problem from real-world safety risk, while still allowing the trained model to be deployed and tested in closed loop — not just evaluated offline — which this project also does (Section 5).

2. Data

The dataset used is [andy8744/udacity-self-driving-car-behavioural-cloning](#) on Kaggle, a repackaging of recordings from the Udacity Self-Driving Car Simulator in the same `driving_log.csv` / `IMG` format used by Udacity's own Self-Driving Car Engineer coursework. It was chosen specifically because it matches a well-documented reference architecture and evaluation methodology (PilotNet), and — critically — because the same simulator that generated the data is itself downloadable and drivable, which let this project validate the trained model in closed loop rather than relying on offline metrics alone.

The dataset contains two separate recordings, corresponding to two visually distinct tracks: a lake-side track (3,930 rows) and a jungle-themed track (3,404 rows), for 7,334 rows total. Each row represents one

timestep and includes three synchronized camera images (center, left, right offset by a fixed lateral distance) plus the human operator's recorded steering angle at that instant (throttle, brake, and speed are also recorded but unused in this project).




Figure 1. Steering angle distribution across both tracks combined. The dataset is heavily skewed toward near-zero (straight-line) steering, since most driving on either track is straight-line road-following. This imbalance directly motivates the straight-frame downsampling and left/right balancing used in the optimized configuration (Section 4).

Size and quality. At $\sim 7,334$ unique timesteps (effectively $\sim 22,000$ images across all three cameras), this is a small dataset by deep-learning standards. It was collected by a single human operator driving with keyboard/mouse input, which introduces some recording noise (small involuntary steering corrections) but is otherwise clean — there are no missing images, corrupted files, or inconsistent formats once the per-track path resolution described in Section 3 is applied.

Expected limitations. Four limitations were expected and confirmed: (1) the data is Unity-rendered, so its visual statistics (flatter shading, repeating textures, no sensor noise) differ from real-world imagery; (2) the dataset is small for a vision task, with no official train/validation split; (3) it reflects one human's driving style on exactly two tracks, with no adverse conditions (rain, night, traffic) represented; and (4) the domain gap in (1) proved empirically significant, not just theoretical — a side investigation (Section 6) found a ResNet18 backbone pretrained on real ImageNet photographs transferred meaningfully *worse* than the compact from-scratch PilotNet used here, evidence that generic pretrained features are a weaker prior on Unity-rendered imagery than on real photographs.

3. Model Setup: Architecture and Preprocessing

Architecture. The model is PilotNet as described by Bojarski et al.: five convolutional layers ($24 \rightarrow 36 \rightarrow 48 \rightarrow 64 \rightarrow 64$ channels; the first three use 5×5 kernels with stride 2, the last two use 3×3 kernels with stride 1), each followed by batch normalization and a ReLU activation, flattened into a four-layer fully-connected regression head ($1152 \rightarrow 126 \rightarrow 64 \rightarrow 32 \rightarrow 1$) with dropout ($p=0.5$) between the first three fully-connected layers. The output is a single unbounded scalar (the predicted steering angle), with no output activation, since this is a regression task. The full model has 287,339 trainable parameters (131,820 in the convolutional backbone, 155,519 in the fully-connected head).

This architecture was chosen for three reasons: it is small enough to train from scratch on CPU-only hardware (no GPU was available) in a practical amount of time; it is the reference architecture specifically documented for this exact task, an externally-validated starting point rather than an arbitrary design; and, somewhat unexpectedly, a side-by-side comparison against a much larger pretrained ResNet18 backbone

(11.2M parameters) found the compact from-scratch PilotNet matched or exceeded it on this data (Section 6), directly supporting the architecture choice here.

Preprocessing pipeline

Each frame is processed as follows before being fed to the network:

1. **Crop** the top 60 pixels (sky) and bottom 25 pixels (the car's hood), neither of which carries steering-relevant information; this reduces irrelevant input variance and focuses the network's receptive field on the road.
2. **Resize** to 200×66 pixels, matching the PilotNet paper's input resolution and keeping the network small and fast to train.
3. **Convert color space from BGR to YUV**, following the reference architecture exactly; this choice was inherited from the paper rather than re-derived, and was not varied in this project.
4. **Normalize** pixel values to approximately $[-1, 1]$ via $x/127.5 - 1$.
5. **Left/right camera substitution** with a fixed ± 0.2 steering correction: since each timestep includes three simultaneous camera views, the left/right images can be treated as "what the center camera would see if the car were off-center," with the steering target adjusted accordingly. This both triples the effective training signal and, more importantly, teaches the model recovery behavior — steering back toward the lane center — which a model trained only on center-camera frames would never see.
6. **Horizontal flip** (50% probability, steering angle negated), a free and exact augmentation exploiting the left-right symmetry of driving.

Regularization. Batch normalization after every convolutional layer and dropout in the fully-connected head were added after early experiments (without them) showed validation loss diverging from training loss — a clear overfitting signature at this data scale. Weight decay (L2 regularization, $1e-4$) and an adaptive learning-rate schedule (`ReduceLROnPlateau`) were added for the same reason and are part of the "optimized" configuration described next.

4. Baseline vs. Optimized Configuration

To satisfy the requirement of comparing against a simpler baseline, two configurations of the *same architecture*, trained and evaluated on the *identical* 80/20 train/validation split of the combined two-track dataset, are compared. Holding architecture and data fixed isolates the effect of the training-recipe changes below.

Table 1. Baseline vs. optimized training configuration.

Aspect	Baseline	Optimized
Architecture	PilotNet (identical in both)	
Preprocessing	Crop / resize / YUV / normalize (identical in both)	
Camera + flip augmentation	Yes (identical in both)	
Class rebalancing	None — plain shuffled mini-batches over the natural ~62:38 straight:turning split (Figure 1)	Straight frames physically downsampled to 60% of their split count (2,165 of 3,609); left-turn frames kept 100% real (1,368); right-turn frames padded with synthetic flipped-left-camera mirrors to match the left-turn count exactly — still a plain shuffled <code>DataLoader</code> , no reweighting sampler
Learning rate	Fixed (Adam, $1e-3$)	Adaptive (<code>ReduceLROnPlateau</code>) starting from $1e-3$

Weight decay	None	1e-4 (L2)
Extra augmentation	None	Random horizontal translation ($\pm 25\text{px}$, steering corrected by $0.004/\text{px}$) + brightness jitter
Loss function	Mean squared error (identical in both)	

The class-rebalancing mechanism went through several iterations, and the search itself is one of this project's more informative findings. An initial `WeightedRandomSampler` (reweighting per-epoch sampling probability, not the underlying data) fixed the straight:turning balance but left a systematic left-steering bias uncorrected, traced to a genuine $\sim 60:40$ left:right skew in the raw turning frames (1,368 left vs. 890 right). A symmetric flip-doubling fix — mirroring every turning frame into the opposite class — corrected left:right balance but diluted both classes equally with synthetic data ($R^2=0.473$). Physically downsampling straight frames was then swept across three keep-rates: 15% regressed sharply (a too-small, low-diversity straight pool produced a *new* left bias, only 17.7% of straight frames correctly classified); 60% modestly beat the sampler baseline; 50% was a small step back from 60%. Finally, replacing symmetric flip-doubling with the asymmetric scheme above — left 100% real, only right padded to parity — gave the best result of the project, improving R^2 , direction accuracy, and large-error rate together rather than trading one for another (Section 5). The lesson: *which* rows are real vs. synthetic padding mattered more than simply equalizing class counts.

Augmentation intensity was tuned separately and empirically: intensities reported in prior community implementations of this task ($\pm 50\text{px}$ translation, $0.4\text{--}1.3\times$ brightness) destabilized training for this smaller, from-scratch model, so intensity was reduced to $\pm 25\text{px}$ / $0.6\text{--}1.2\times$, discussed further as a failure case in Section 5.

5. Evaluation

Six metrics are reported. **RMSE** and **MAE** are in the same units as the steering angle itself (roughly $[-1, 1]$) and are directly interpretable. **R^2** compares the model's mean squared error against a trivial baseline that always predicts the mean steering angle; this metric matters specifically because of the class imbalance shown in Figure 1 — a model that has simply learned to predict approximately zero for every frame can post a deceptively low raw MSE, since roughly 62% of validation frames genuinely have near-zero steering. R^2 makes this failure mode visible. Because straight-line frames dominate the dataset and are easy (predicting ~ 0 is nearly correct), **RMSE is additionally reported separately for straight ($|\text{angle}| < 0.05$) and turning ($|\text{angle}| \geq 0.05$) frames**, since an aggregate average can hide poor turning performance behind a large majority of easy straight frames.

RMSE-based metrics share a blind spot: they penalize any magnitude difference equally, even when the model chose the correct steering *direction* and simply differed on how sharply to turn — not necessarily a real mistake, given this data reflects one human's driving style rather than a single correct answer. Two further metrics target the more driving-relevant question of whether the model made the correct qualitative decision: **direction accuracy** (predicted and actual steering agree on straight/left/right, threshold 0.05) and **large-error rate** ($|\text{actual} - \text{predicted}| > 0.3$, a proxy for outright dangerous mispredictions). A 5-bucket confusion matrix (hard-left/soft-left/straight/soft-right/ hard-right, hard cutoff 0.25) makes the specific failure pattern visible (Figure 2, bottom).

Table 2. Baseline vs. optimized evaluation metrics on the shared 1,200-frame validation set.

Metric	Baseline	Optimized
MSE	0.06036	0.05476
RMSE	0.24569	0.23401
MAE	0.16978	0.16361

R^2	0.437	0.489
Straight-frame RMSE (n=747)	0.13031	0.14038
Turning-frame RMSE (n=453)	0.36318	0.33550
Direction accuracy (3-way)	51.2%	54.6%
Large-error rate ($ \text{error} > 0.3$)	15.1%	13.5%




Figure 2. Top: baseline vs. optimized configuration across error metrics (left) and R^2 /direction accuracy/large-error rate (right) — every metric moves in the optimized configuration's favor. Bottom: 5-bucket steering confusion matrices on the shared 1,200-frame validation set (rows=actual, cols=predicted, shaded by row %). The dominant error in both configurations is straight frames predicted as soft-left — see discussion below.

Discussion. The final configuration shows a real, consistent gap over baseline across every metric at once: R^2 improves by 0.052 (0.437→0.489), MSE drops 9.3%, direction accuracy improves 3.4 points (51.2%→54.6%), large-error rate falls 1.6 points (15.1%→13.5%). Re-running identical code at earlier stages of this project produced R^2 swings of several hundredths from random initialization alone, so no single metric here is strong evidence in isolation — but four largely independent measures (R^2 , direction accuracy, large-error rate, and the turning/straight RMSE split) all moving together at once is much less likely by chance. As before, there is a genuine trade: straight-frame RMSE gets slightly worse (0.130→0.140) while turning-frame RMSE improves (0.363→0.336), the same operationally-favorable trade as before, now larger in magnitude.

Failure cases

- Turning remains the dominant error source in both configurations** — turning-frame RMSE is 2.4–2.8× straight-frame RMSE in both configurations (0.363 vs. 0.130 baseline; 0.336 vs. 0.140 optimized), a ratio that held across essentially every configuration tested in this project, including alternative architectures explored outside this report's main comparison. The model is disproportionately worse exactly where accuracy matters most for staying on the road.
- A persistent left-leaning bias on straight road.** Figure 2 (bottom) shows the largest single source of misclassification in both configurations is straight frames predicted as soft-left — 312/747 baseline, 275/747 optimized. Present even in the baseline (which uses none of Section 4's left/right

balancing), this points to a property of the raw data or architecture, not something rebalancing introduced. Targeted left/right-balancing experiments narrowed it (right-turn accuracy 88.2%→96.1%) but did not eliminate it — the clearest specific weakness of the final model.

3. **Open-loop metrics do not fully predict closed-loop behavior.** An earlier optimized configuration was deployed live in the Udacity simulator via a small Socket.IO bridge script and completed extended stretches of both tracks autonomously. This is not something the offline metrics above measure directly: they evaluate one frame at a time against a fixed, human-driven trajectory (open-loop), whereas driving is closed-loop — small per-frame errors compound as the car drifts into road positions under-represented in training data, a well-known distributional-shift problem in imitation learning. The final configuration above was not independently re-verified live.
4. **Augmentation intensity does not transfer directly from other work.** Intensities standard in prior community implementations of this task measurably destabilized training here (Section 4); testing three levels found a monotonic negative relationship between intensity and validation performance for this architecture/data scale — augmentation hyperparameters must be re-validated per model/dataset rather than assumed to transfer.

6. Interpretation

What did the model learn? The model receives only raw pixels and a scalar steering target — no lane-line, object, or semantic labels — so any competence must come from implicitly associating road-edge and lane-curvature cues with the correct steering response. Several design choices only make sense, and only work, if this is what the network learned: the crop step (Section 3) restricts input to the road region; the left/right-camera augmentation only teaches useful "recovery" behavior if the network uses lateral road-position information rather than memorizing textures; and flip augmentation only helps if the network's features are genuinely spatial/geometric rather than track-specific memorization. That flip augmentation measurably helps (part of every configuration including baseline; removing it during early experiments produced visibly worse turning behavior) is itself indirect evidence for "lane geometry" over memorization.

Where can it fail? Four failure modes were observed directly in this project: sharp or unusual turns (the turning-RMSE gap, Section 5); a left-leaning bias on genuinely straight road that persisted across every configuration tested including the baseline, and which targeted rebalancing narrowed but did not remove (Section 5), suggesting a root cause in the raw data or architecture rather than class-count imbalance alone; track-specific visual style, where brightness-jitter augmentation aimed at generalizing across the lake and jungle tracks' different color palettes produced at best a neutral effect — two tracks may simply be too few to learn a genuinely lighting-invariant representation; and, more fundamentally, that the model can only handle situations represented in its training demonstrations at all — a structural limitation of behavioral cloning itself, and the reason the closed-loop distributional-shift issue in Section 5 exists in the first place.

A domain-gap finding. A side investigation trained a ResNet18 backbone pretrained on ImageNet, first with only its final residual block unfrozen and later fully fine-tuned, against the PilotNet models above. Fully frozen (a small regression head on fixed ImageNet features), it collapsed into predicting near-zero steering for almost every frame — the *best* straight-frame RMSE of any configuration tested, but by a wide margin the *worst* turning-frame RMSE, the clearest signature of a model minimizing error on the easy majority class rather than tracking the road. Unfreezing more of the backbone improved this, but even a fully fine-tuned ResNet18 did not clearly exceed the compact from-scratch PilotNet. This suggests "what a model learns" depends as much on whether its input representation matches the task's visual domain as on raw capacity — more parameters is not automatically better when source and target domains (real photographs vs. a Unity-rendered simulator) differ this much.

7. Conclusion: Why This Data and Setup Were the Right Choice (or Not)

This dataset and setup were, on balance, the right choice for the goals of this project. Using data recorded from an actual drivable simulator — rather than a static image dataset — made it possible to go beyond offline metrics to a genuine closed-loop test, directly surfacing the open-loop/closed-loop gap discussed in Section 5, a more sophisticated and honest evaluation than most small-scale behavioral-cloning projects attempt. The compact, from-scratch PilotNet architecture was well matched to the CPU-only compute budget available, and the comparison against a pretrained ResNet18 (Section 6) confirms this was not a compromise on quality: the larger, pretrained alternative did not outperform it here.

The setup was not, however, unambiguously right for every goal. The dataset's small size and narrow diversity (two tracks, one driver, no adverse conditions) still appear to place a real ceiling on results: across every training-recipe variation tried, validation R^2 ranged from roughly 0.43 to 0.49, and turning-frame error never closed most of the gap to straight-frame error. Notably, the single largest improvement (R^2 0.437→0.489) came not from a broader recipe change but specifically from *how* class rebalancing was done — real vs. synthetic data treated asymmetrically rather than just equalizing class counts (Section 4) — and a left-leaning bias on straight road resisted direct correction despite several dedicated attempts. This is more consistent with a ceiling on genuinely diverse real data per class than a single implementation flaw: recipe choices moved results meaningfully but could not manufacture diversity the raw recordings don't contain. A larger, more diverse dataset would likely still be a more effective investment than further tuning on the current data.

Summary. For building and deploying a complete behavioral-cloning pipeline end to end — including a genuine closed-loop test, not just offline metrics — this dataset and setup were well-matched. For maximizing absolute accuracy, *how* the available data is used mattered more than architecture or hyperparameter tuning, but the amount of genuinely diverse real data still appears to be the binding constraint on how far that can go.

References

Bojarski, M., Del Testa, D., Dworakowski, D., Firner, B., Flepp, B., Goyal, P., Jackel, L. D., Monfort, M., Muller, U., Zhang, J., Zhang, X., Zhao, J., & Zieba, K. (2016). *End to End Learning for Self-Driving Cars*. NVIDIA Corporation. arXiv:1604.07316.

Kaggle dataset: *Udacity Self Driving Car — Behavioural Cloning*, andy8744/udacity-self-driving-car-behavioural-cloning.

Udacity, Inc. *Self-Driving Car Simulator*.

EXAM 3 — THEORETICAL QUESTIONS (SEPARATE FROM THE PROJECT TASK REPORT ABOVE)

Theoretical Question 1: Linear Models vs. CNNs vs. RNN/Transformers

Compare regarding typical data types, strengths/weaknesses, and overfitting risk & regularization; one use case each.

Table 3. Linear models vs. CNNs vs. RNN/Transformers.

Aspect	Linear models	CNNs	RNN / Transformers
Typical data	Tabular/structured numeric features	Grid data, local spatial correlation (images, spectrograms)	Sequential/ordered data (text, time series, audio)
Strengths	Interpretable coefficients; fast/convex	Weight sharing captures translation-invariant local patterns; hierarchical	Model order/temporal dependencies; Transformers

	to fit; cheap, strong baseline	features (edges→textures→objects) with few weights vs. a comparable dense net	capture long-range dependencies in parallel via self-attention, avoiding RNNs' sequential bottleneck
Weaknesses	Assumes linear/hand-engineered input→target relation; no automatic hierarchical structure	Limited native long-range context without deep stacking/attention; needs sizeable labeled data or pretraining	RNNs: sequential, slow, vanishing/exploding gradients on long sequences. Transformers: data/compute-hungry, attention cost grows quadratically in sequence length
Overfitting & reg.	Low risk (few params), but occurs with many/collinear features; L1 (Lasso) / L2 (Ridge) penalties	Moderate–high on small data; augmentation, dropout, batch norm, weight decay — the combination this project's own PilotNet uses (Section 3)	High (large param counts); dropout, layer norm, weight decay, early stopping, large-scale pretraining/fine-tuning
Use case	Credit-scoring from applicant features, where interpretability is a regulatory need	This project: steering-angle prediction from dashcam images (PilotNet)	Machine translation (Transformers); sensor/price time-series forecasting (RNN/LSTM)

Theoretical Question 2: Training Neural Networks with Gradient Descent

- 1. Loss function and its gradient.** The loss quantifies how wrong predictions are relative to targets (e.g. MSE in this project's training loop, Section 3). Its gradient w.r.t. each parameter, from backpropagation, points toward *steepest increase* in loss; stepping in the *negative* gradient direction is what reduces loss each update.
- 2. Batch vs. mini-batch vs. stochastic GD.** Batch GD uses the whole training set per update: stable but slow/memory-heavy. Stochastic GD updates per single example: fast per step, and noise can escape shallow local minima, but high variance can prevent settling at an optimum. Mini-batch GD (used here, `BATCH_SIZE=64`) is the practical middle ground almost all deep learning uses, balancing gradient-estimate stability against compute cost and hardware parallelism.
- 3. The learning rate** scales the update step along the negative gradient. Too large: updates overshoot, causing oscillation or divergence — seen directly here when an over-aggressive augmentation/scheduler combination destabilized training (Section 4). Too small: convergence is impractically slow or stalls on plateaus. This project's `ReduceLROnPlateau` addresses both, shrinking the rate only once validation loss stops improving.
- 4. Gradient descent vs. ascent.** Descent moves along the *negative* gradient to *minimize* an objective — the default for supervised learning, including this project's regression. Ascent moves along the *positive* gradient to *maximize* an objective, used for a reward or likelihood to maximize — e.g. policy-gradient RL, or log-likelihood maximization (often reframed as minimizing negative log-likelihood, why descent is more commonly implemented).